Lecture 17: NumPy and StatsModels

Analysis and model fitting in Python

First, the libraries

```
1  import requests
2
3  import numpy as np
4  import statsmodels.api as sm
5  import statsmodels.formula.api as smf
```

Wrap-up from lecture 16: Post Requests

https://ww2.energy.ca.gov/almanac/renewables_data/wind/index_cms.php

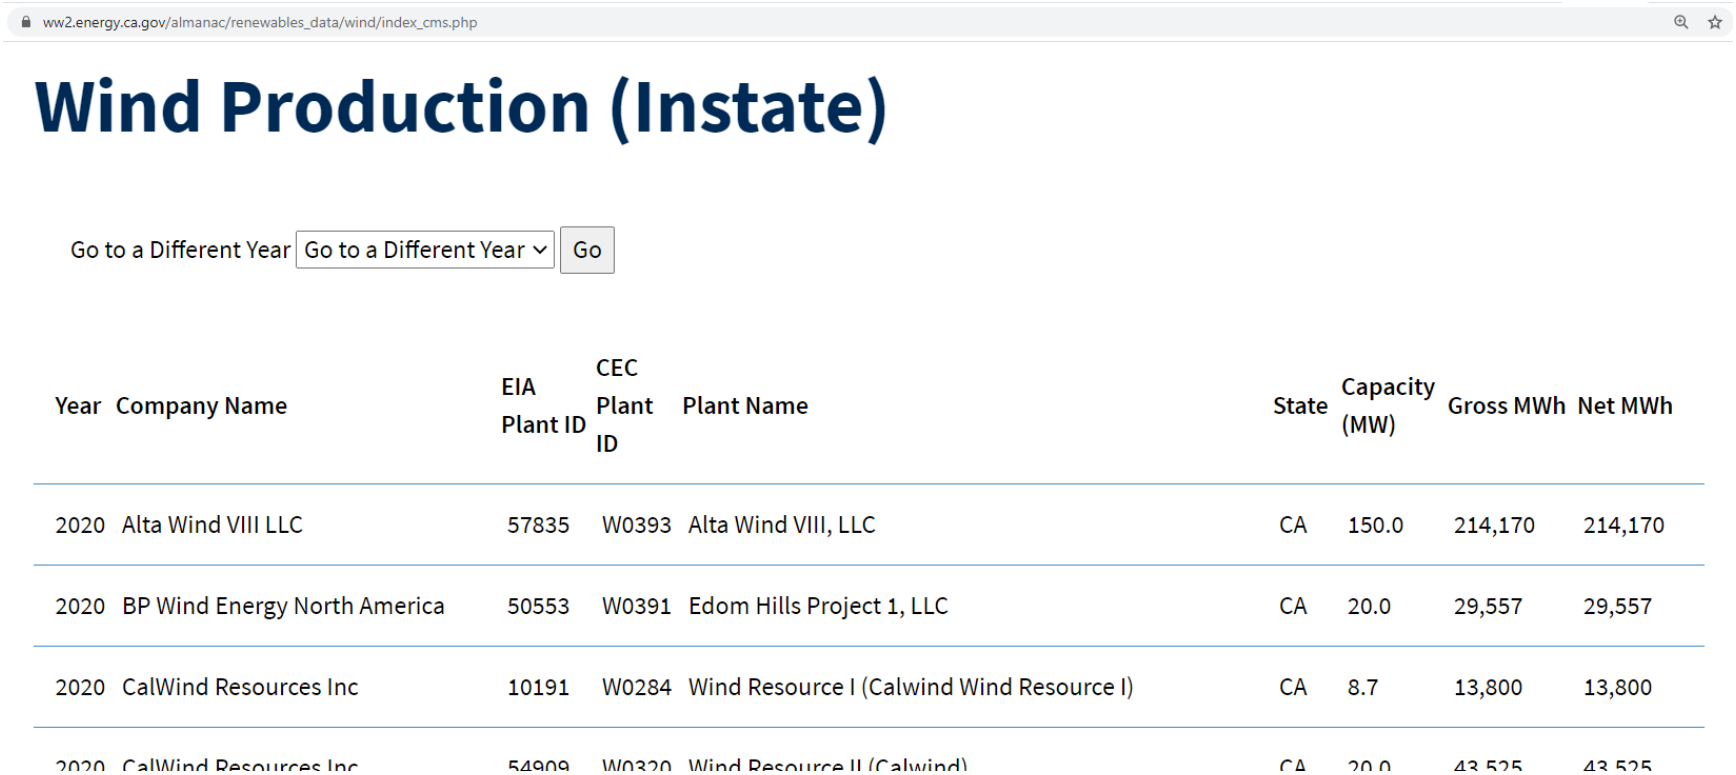

Wrap-up from lecture 16: Post Requests

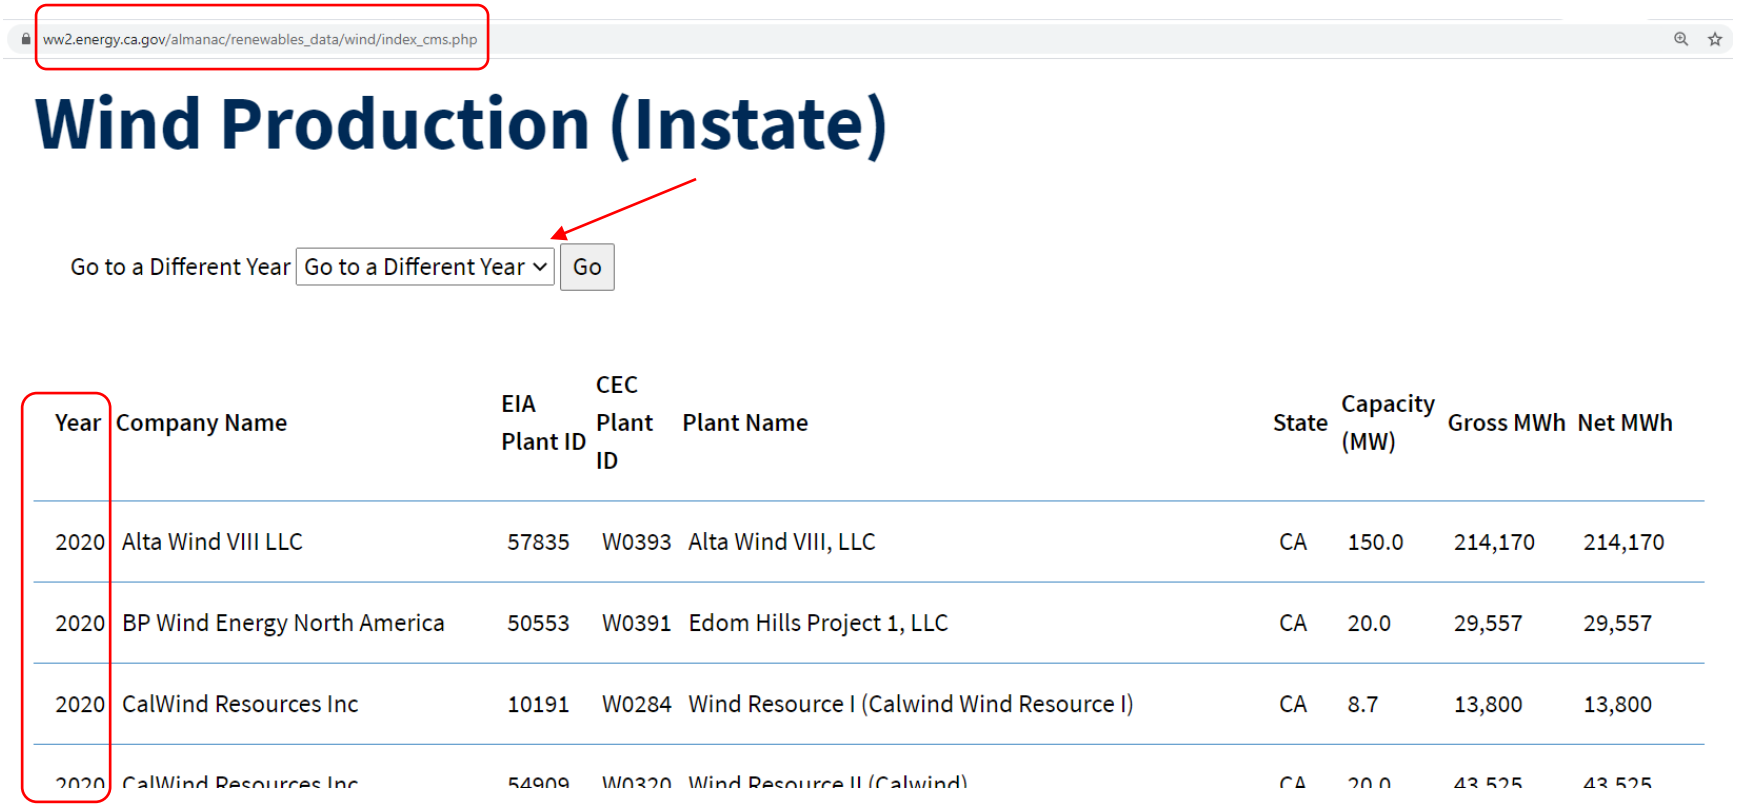

Year	Company Name	EIA Plant ID	CEC Plant ID	Plant Name	State	Capacity (MW)	Gross MWh	Net MWh
2020	Alta Wind VIII LLC	57835	W0393	Alta Wind VIII, LLC	CA	150.0	214,170	214,170
2020	BP Wind Energy North America	50553	W0391	Edom Hills Project 1, LLC	CA	20.0	29,557	29,557
2020	CalWind Resources Inc	10191	W0284	Wind Resource I (Calwind Wind Resource I)	CA	8.7	13,800	13,800
2020	CalWind Resources Inc	54909	W0320	Wind Resource II (Calwind)	CA	20.0	43,525	43,525

Wrap-up from lecture 16: Post Requests

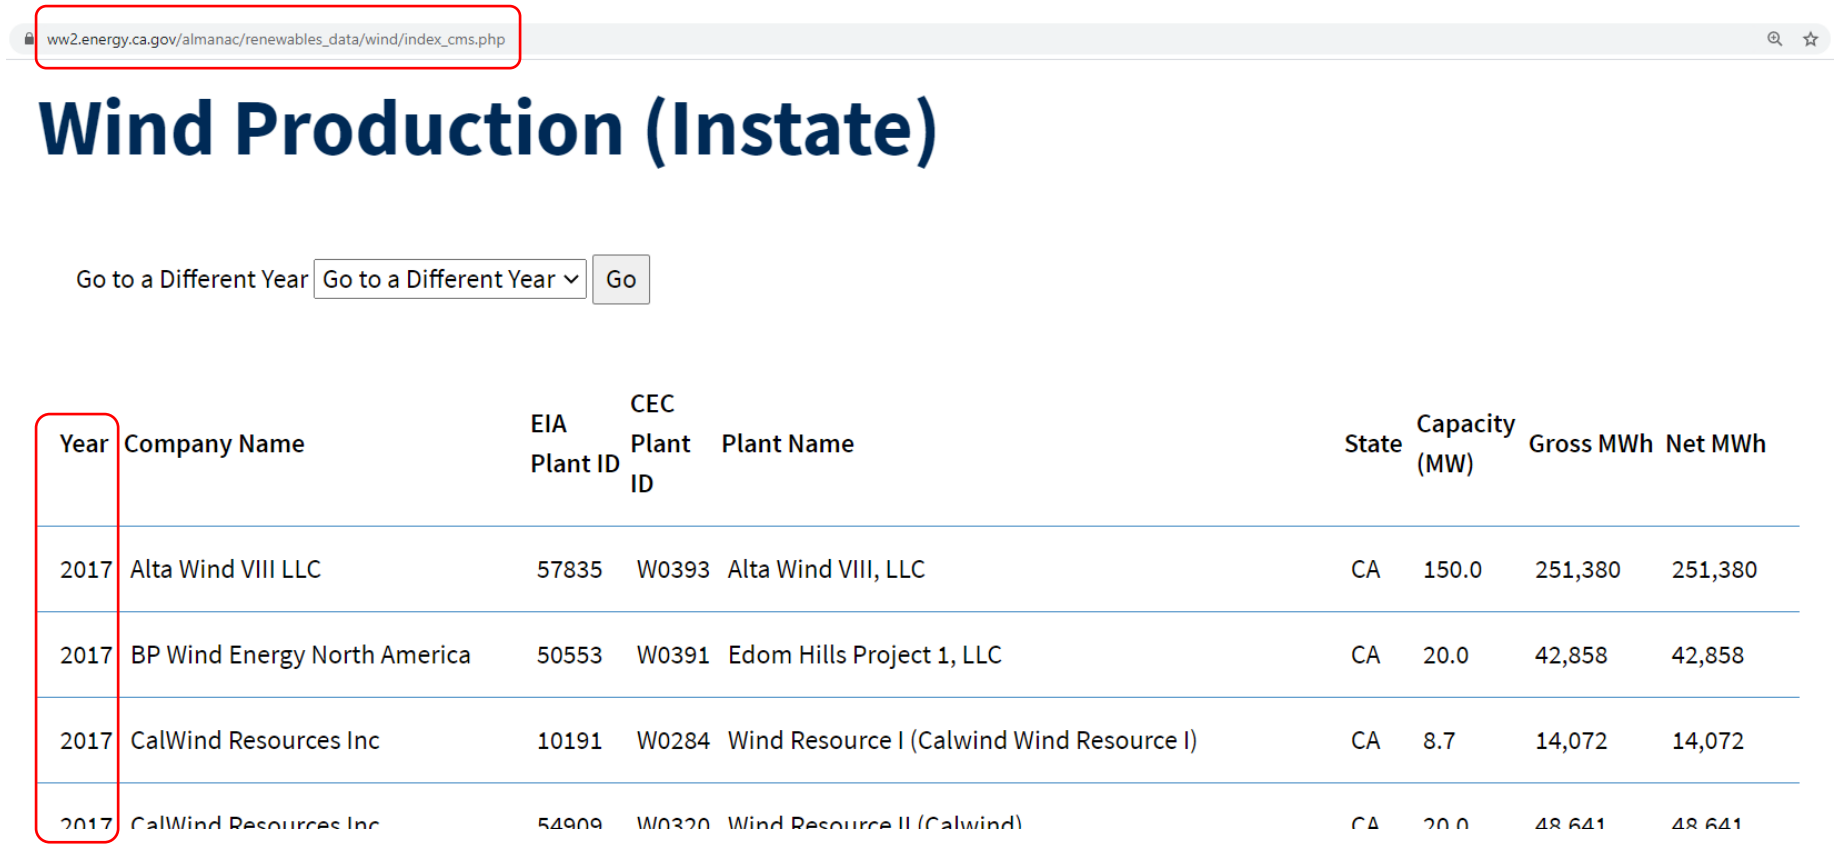

ww2.energy.ca.gov/almanac/renewables_data/wind/index_cms.php

Wind Production (Instate)

Go to a Different Year | Go to a Different Year | Go

Year	Company Name	EIA Plant ID	CEC Plant ID	Plant Name	State	Capacity (MW)	Gross MWh	Net MWh
2017	Alta Wind VIII LLC	57835	W0393	Alta Wind VIII, LLC	CA	150.0	251,380	251,380
2017	BP Wind Energy North America	50553	W0391	Edom Hills Project 1, LLC	CA	20.0	42,858	42,858
2017	CalWind Resources Inc	10191	W0284	Wind Resource I (Calwind Wind Resource I)	CA	8.7	14,072	14,072
2017	CalWind Resources Inc	54909	W0320	Wind Resource II (Calwind)	CA	20.0	48,641	48,641

Wrap-up from lecture 16: Post Requests

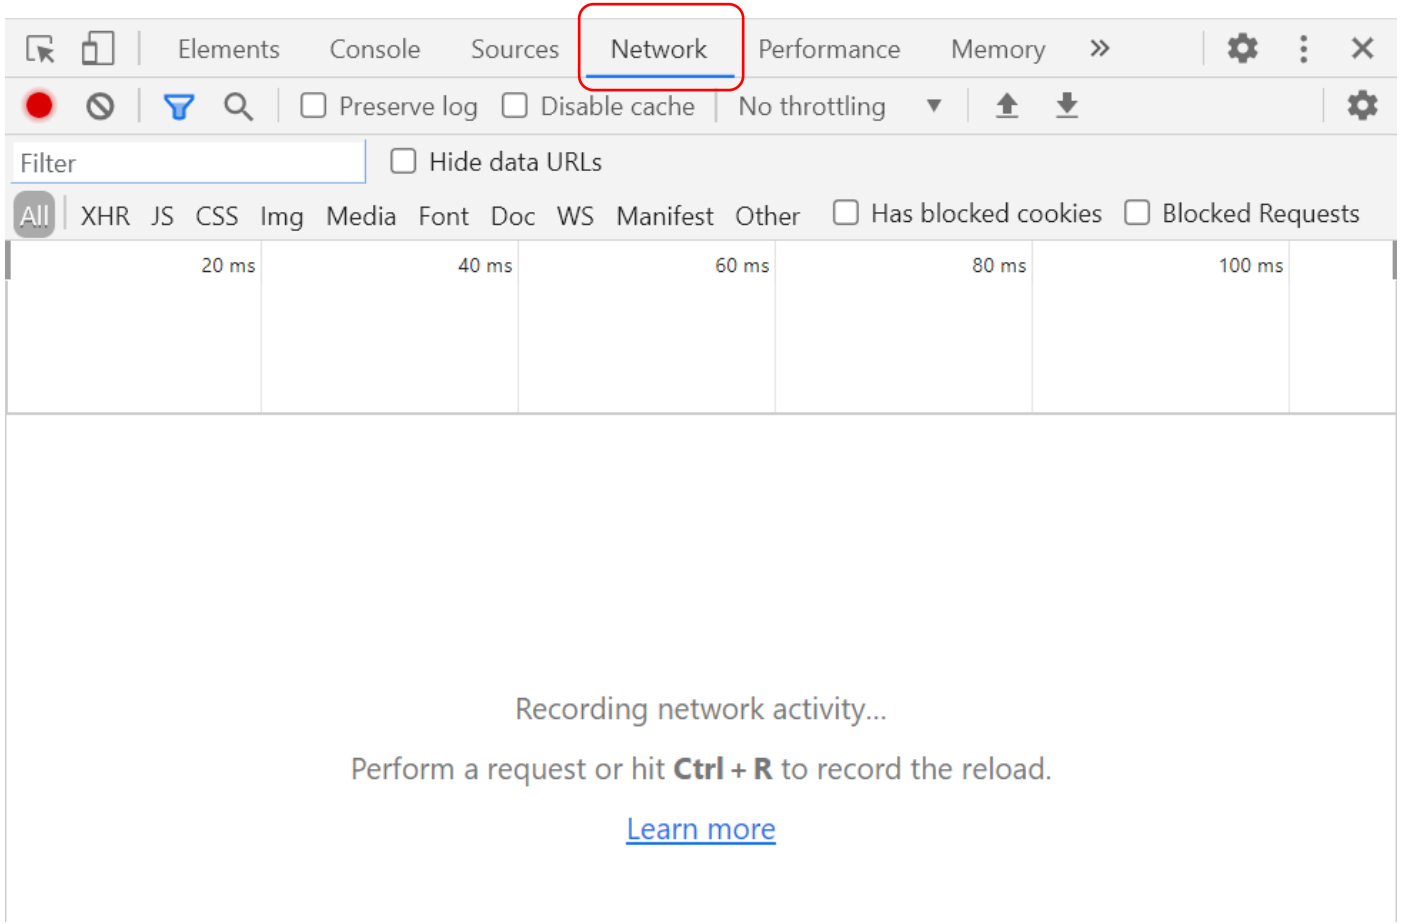

Wrap-up from lecture 16: Post Requests

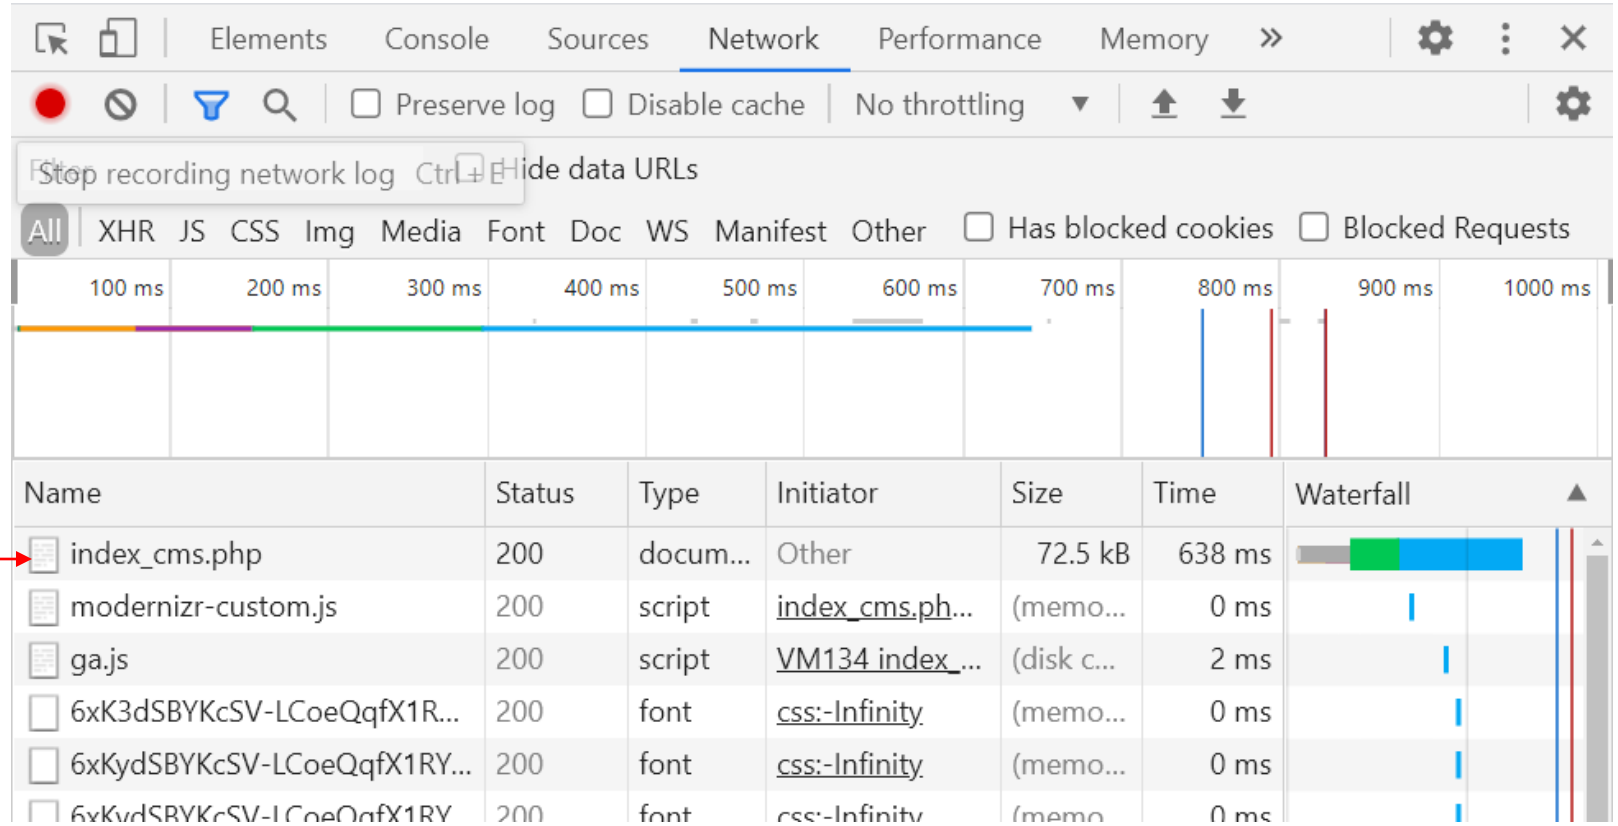

Wrap-up from lecture 16: Post Requests

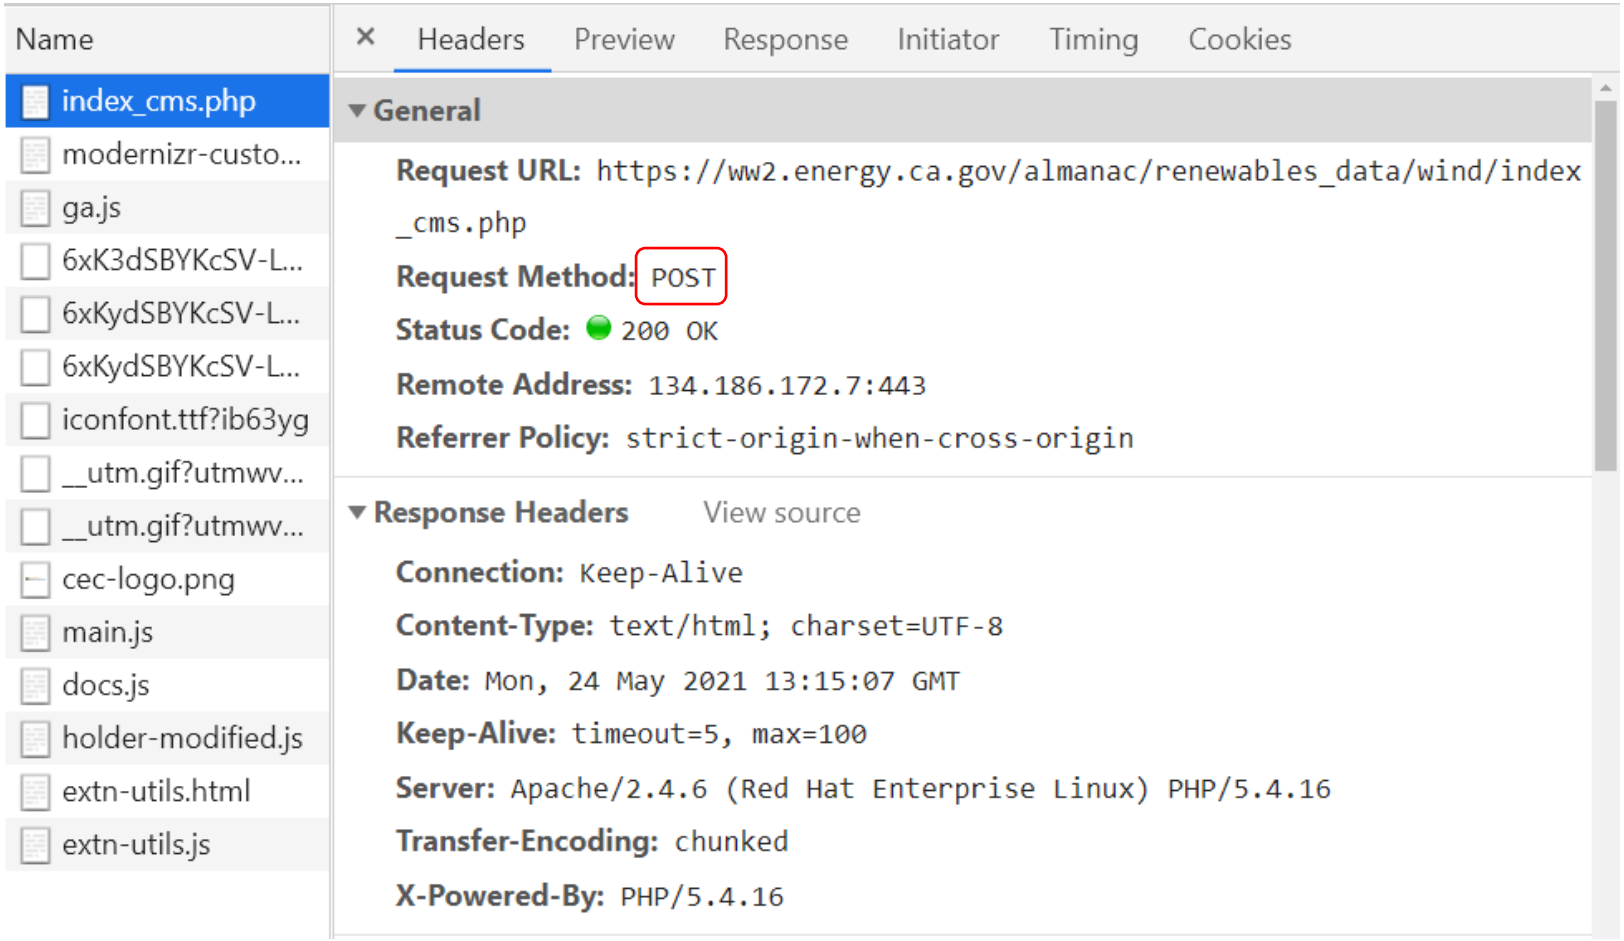

Wrap-up from lecture 16: Post Requests

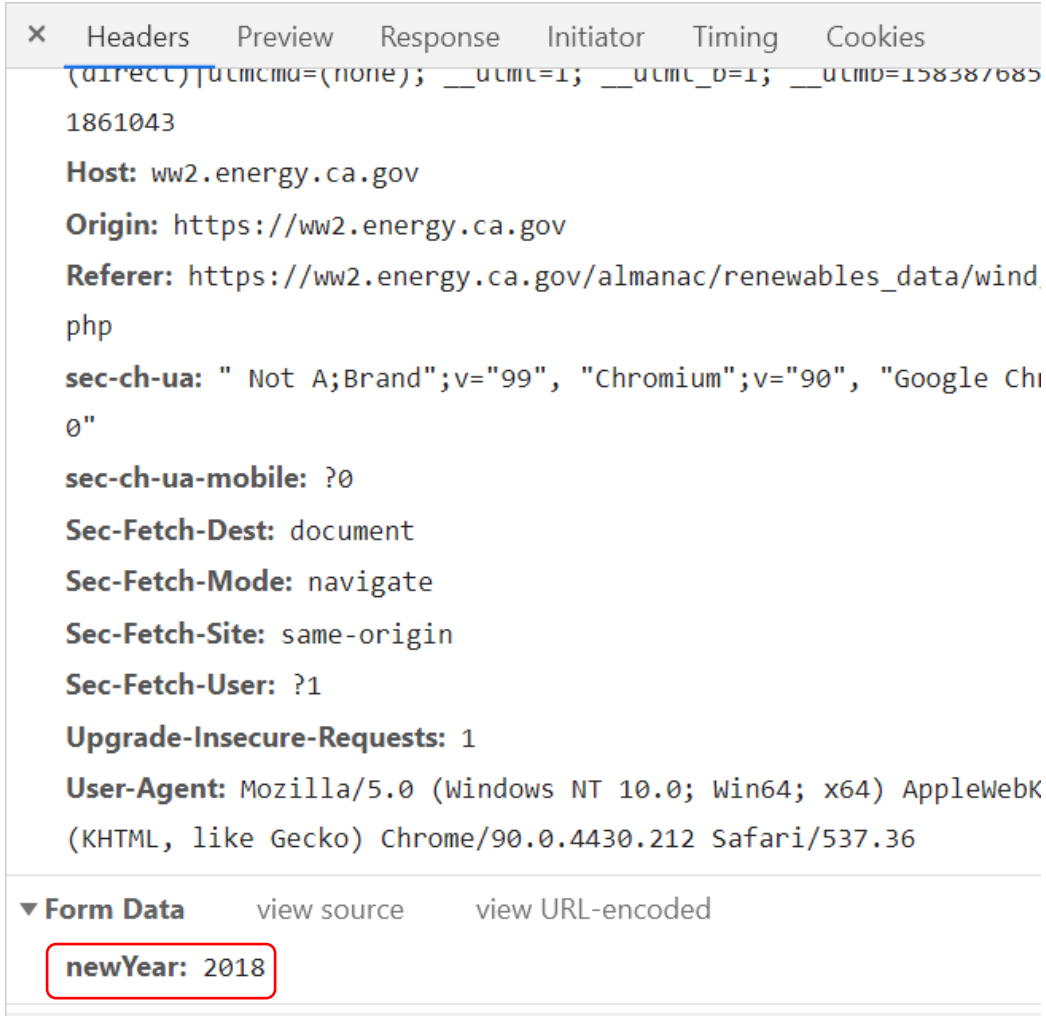

Wrap-up from lecture 16: Post Requests

```
8   response = requests.post(url)
9
```

```
In [2]: response.text.find('Alta Wind VIII LLC')
Out[2]: 34328
```

```
In [4]: response.text[34300:34600]
Out[4]: "tr>\n\t\t\n<tr><td>2020</td><td>Alta Wind VIII
LLC</td><td>57835</td><td>W0393</td><td>Alta Wind VIII,
LLC</td><td>CA</td><td class='right'>150.0</td><td
class='right'>214,170</td><td class='right'>214,170</
td></tr>\r\t\n<tr><td>2020</td><td>BP Wind Energy North
America</td><td>50553</td><td>W0391</td><td>"
```

Wrap-up from lecture 16: Post Requests

```
8   response = requests.post(url)
9
```

```
In [2]: response.text.find('Alta Wind VIII LLC')
Out[2]: 34328
```

```
In [4]: response.text[34300:34600]
Out[4]: "tr>\n\t\t\n<tr><td>2020</td><td>Alta Wind VIII
LLC</td><td>57835</td><td>W0393</td><td>Alta Wind VIII,
LLC</td><td>CA</td><td class='right'>150.0</td><td
class='right'>214,170</td><td class='right'>214,170</
td></tr>\r\t\n<tr><td>2020</td><td>BP Wind Energy North
America</td><td>50553</td><td>W0391</td><td>"
```

Wrap-up from lecture 16: Post Requests

```
8   response = requests.post(url)
9
```

```
In [2]: response.text.find('Alta Wind VIII LLC')
Out[2]: 34328
```

```
In [4]: response.text[34300:34600]
Out[4]: "tr>\n\t\t\n<tr><td>2020</td><td>Alta Wind VIII
LLC</td><td>57835</td><td>W0393</td><td>Alta Wind VIII,
LLC</td><td>CA</td><td class='right'>150.0</td><td
class='right'>214,170</td><td class='right'>214,170</
td></tr>\r\t\n<tr><td>2020</td><td>BP Wind Energy North
America</td><td>50553</td><td>W0391</td><td>"
```

Year	Company Name	EIA Plant ID	Plant ID	Plant Name	State	Capacity (MW)	Gross MWh	Net MWh
2020	Alta Wind VIII LLC	57835	W0393	Alta Wind VIII, LLC	CA	150.0	214,170	214,170

Wrap-up from lecture 16: Post Requests

```
13    year = 2017
14    response = requests.post(url, data={'newYear':year})
```

```
16    response.text.find('Alta Wind VIII LLC')
```

```
In [11]: response.text[34300:34550]
Out[11]: "tr>\n\t\t\n<tr><td>2017</td><td>Alta Wind VIII
LLC</td><td>57835</td><td>W0393</td><td>Alta Wind VIII,
LLC</td><td>CA</td><td class='right'>150.0</td><td
class='right'>251,380</td><td class='right'>251,380</
td></tr>\r\t\n<tr><td>2017</td><td>BP Wind Energy "
```

Wrap-up from lecture 16: Post Requests

```
13  year = 2017
14  response = requests.post(url, data={'newYear':year})

16  response.text.find('Alta Wind VIII LLC')
```

```
In [11]: response.text[34300:34550]
Out[11]: "tr>\n\t\t\n<tr><td>2017</td><td>Alta Wind VIII
LLC</td><td>57835</td><td>W0393</td><td>Alta Wind VIII,
LLC</td><td>CA</td><td class='right'>150.0</td><td
class='right'>251,380</td><td class='right'>251,380</
td></tr>\r\t\n<tr><td>2017</td><td>BP Wind Energy "
```

Year	Company Name	EIA Plant ID	Plant ID	Plant Name	State	Capacity (MW)	Gross MWh	Net MWh
2017	Alta Wind VIII LLC	57835	W0393	Alta Wind VIII, LLC	CA	150.0	251,380	251,380

Numpy

```
20  my_list  = [1, 2, 3, 4, 5]
21  my_array = np.array([1, 2, 3, 4, 5])
```

Numpy

```
20    my_list  = [1, 2, 3, 4, 5]
21    my_array = np.array([1, 2, 3, 4, 5])
```

```
In [24]: my_list*3
Out[24]: [1, 2, 3, 4, 5, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5]

In [25]: my_array*3
Out[25]: array([ 3,  6,  9, 12, 15])
```

```
In [26]: my_list + my_list
Out[26]: [1, 2, 3, 4, 5, 1, 2, 3, 4, 5]

In [27]: my_array + my_array
Out[27]: array([ 2,  4,  6,  8, 10])
```

Numpy

```
20    my_list  = [1, 2, 3, 4, 5]
21    my_array = np.array([1, 2, 3, 4, 5])
```

```
In [28]: np.concatenate((my_array, my_array))
Out[28]: array([1, 2, 3, 4, 5, 1, 2, 3, 4, 5])
```

Numpy

```
20    my_list  = [1, 2, 3, 4, 5]
21    my_array = np.array([1, 2, 3, 4, 5])
```

```
In [28]: np.concatenate((my_array, my_array))
Out[28]: array([1, 2, 3, 4, 5, 1, 2, 3, 4, 5])
```

```
In [29]: np.stack((my_array, my_array), axis=0)
Out[29]:
array([[1, 2, 3, 4, 5],
       [1, 2, 3, 4, 5]])
```

Numpy

```
20    my_list  = [1, 2, 3, 4, 5]
21    my_array = np.array([1, 2, 3, 4, 5])
```

```
In [28]: np.concatenate((my_array, my_array))
Out[28]: array([1, 2, 3, 4, 5, 1, 2, 3, 4, 5])
```

```
In [29]: np.stack((my_array, my_array), axis=0)
Out[29]:
array([[1, 2, 3, 4, 5],
       [1, 2, 3, 4, 5]])
```

```
In [30]: my_array.reshape(5, 1)
Out[30]:
array([[1],
       [2],
       [3],
       [4],
       [5]])
```

Numpy

```
In [31]: np.array([1, 2, 'cat'])
Out[31]: array(['1', '2', 'cat'], dtype='<U11')
```

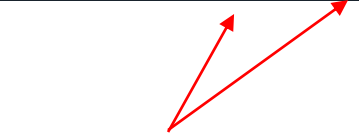

Became strings

Numpy

```
38    mat = np.array([[1, 2], [3, 4], [5, 6]])
```

```
array([[1, 2],
       [3, 4],
       [5, 6]])
```

```
In [34]: mat[1]
Out[34]: array([3, 4])

In [35]: mat[1][1]
Out[35]: 4

In [36]: mat[1, 1]
Out[36]: 4
```

New to Numpy arrays

Numpy: matrix algebra

```
In [39]: mat.T
Out[39]:
array([[1, 3, 5],
       [2, 4, 6]])
```

Numpy: matrix algebra

```
In [39]: mat.T
Out[39]:
array([[1, 3, 5],
       [2, 4, 6]])
```

```
In [40]: mat.dot(mat.T)
Out[40]:
array([[ 5, 11, 17],
       [11, 25, 39],
       [17, 39, 61]])
```

Numpy: matrix algebra

```
In [39]: mat.T
Out[39]:
array([[1, 3, 5],
       [2, 4, 6]])
```

```
In [40]: mat.dot(mat.T)
Out[40]:
array([[ 5, 11, 17],
       [11, 25, 39],
       [17, 39, 61]])
```

```
In [41]: np.linalg.pinv(mat)
Out[41]:
array([[-1.33333333, -0.33333333,  0.66666667],
       [ 1.08333333,  0.33333333, -0.41666667]])
```

Statsmodels

```
53  df = sm.datasets.get_rdataset('Guerry', 'HistData').data
54  df.head()
```

	dept	Region	Department	Crime_pers	...	Prostitutes	Distance	Area	Pop1831
0	1	E	Ain	28870	...	13	218.372	5762	346.03
1	2	N	Aisne	26226	...	327	65.945	7369	513.00
2	3	C	Allier	26747	...	34	161.927	7340	298.26
3	4	E	Basses-Alpes	12935	...	2	351.399	6925	155.90
4	5	E	Hautes-Alpes	17488	...	1	320.280	5549	129.10

[5 rows x 23 columns]

Statsmodels: R-style formulas

```
56  model = smf.ols('Lottery ~ Literacy + np.log(Pop1831)', data=df)
57  result = model.fit()
58  rs = result.summary()
```

```
"""
                            OLS Regression Results
==============================================================================
Dep. Variable:                Lottery   R-squared:                       0.348
Model:                            OLS   Adj. R-squared:                  0.333
Method:                 Least Squares   F-statistic:                     22.20
Date:                Mon, 24 May 2021   Prob (F-statistic):           1.90e-08
Time:                        10:59:47   Log-Likelihood:                -379.82
No. Observations:                  86   AIC:                             765.6
Df Residuals:                      83   BIC:                             773.0
Df Model:                           2
Covariance Type:            nonrobust
==============================================================================
                 coef    std err          t      P>|t|      [0.025      0.975]
------------------------------------------------------------------------------
Intercept      246.4341     35.233      6.995      0.000     176.358     316.510
Literacy        -0.4889      0.128     -3.832      0.000      -0.743      -0.235
np.log(Pop1831) -31.3114      5.977     -5.239      0.000     -43.199     -19.424
==============================================================================
Omnibus:                        3.713   Durbin-Watson:                   2.019
Prob(Omnibus):                  0.156   Jarque-Bera (JB):                3.394
Skew:                          -0.487   Prob(JB):                        0.183
Kurtosis:                       3.003   Cond. No.                         702.
==============================================================================
```

Statsmodels: R-style formulas

```
56  model = smf.ols('Lottery ~ Literacy + np.log(Pop1831)', data=df)
57  result = model.fit()
58  rs = result.summary()
```

```
In [51]: result.pvalues
Out[51]:
Intercept           6.260771e-10
Literacy            2.462102e-04
np.log(Pop1831)     1.202925e-06
dtype: float64

In [52]: result.params
Out[52]:
Intercept           246.434135
Literacy             -0.488923
np.log(Pop1831)     -31.311392
dtype: float64

In [53]: result.rsquared
Out[53]: 0.3484706112599609
```

Statsmodels: matrices

```
In [55]: df.to_numpy()
Out[55]:
array([[1, 'E', 'Ain', ..., 218.372, 5762, 346.03],
       [2, 'N', 'Aisne', ..., 65.945, 7369, 513.0],
       [3, 'C', 'Allier', ..., 161.927, 7340, 298.26],
       ...,
       [88, 'E', 'Vosges', ..., 174.477, 5874, 397.99],
       [89, 'C', 'Yonne', ..., 81.797, 7427, 352.49],
       [200, nan, 'Corse', ..., 539.213, 8680, 195.41]], dtype=object)
```

Statsmodels: matrices

```
65  df['logpop'] = np.log(df['Pop1831'])
66  df['intercept'] = np.ones(len(df))
67  model = sm.OLS(endog=df['Lottery'], exog=df[['intercept', 'Literacy', 'logpop']])
68  result = model.fit()
69  result.summary()
```

```
"""
                            OLS Regression Results
==============================================================================
Dep. Variable:                Lottery   R-squared:                       0.348
Model:                            OLS   Adj. R-squared:                  0.333
Method:                 Least Squares   F-statistic:                     22.20
Date:                Mon, 24 May 2021   Prob (F-statistic):           1.90e-08
Time:                        11:07:12   Log-Likelihood:                -379.82
No. Observations:                  86   AIC:                             765.6
Df Residuals:                      83   BIC:                             773.0
Df Model:                           2
Covariance Type:            nonrobust
==============================================================================
                 coef    std err          t      P>|t|      [0.025      0.975]
------------------------------------------------------------------------------
intercept    246.4341     35.233      6.995      0.000     176.358     316.510
Literacy      -0.4889      0.128     -3.832      0.000      -0.743      -0.235
logpop       -31.3114      5.977     -5.239      0.000     -43.199     -19.424
==============================================================================
Omnibus:                        3.713   Durbin-Watson:                   2.019
Prob(Omnibus):                  0.156   Jarque-Bera (JB):                3.394
Skew:                          -0.487   Prob(JB):                        0.183
Kurtosis:                       3.003   Cond. No.                         702.
==============================================================================
```